

***note:***  
***black & white***




# Recap – Image Classification with Bags of Local Features

- Bag of Feature models were the state of the art for image classification for a decade
- BoF may still be the state of the art for instance retrieval
- We saw numerous strategies to fight back against lost spatial information (spatial pyramid) and lost feature detail due to quantization.
- Food for thought, doesn't the spatial pyramid seem kind of recursive / hierarchical? Like a SIFT feature on top of SIFT features?

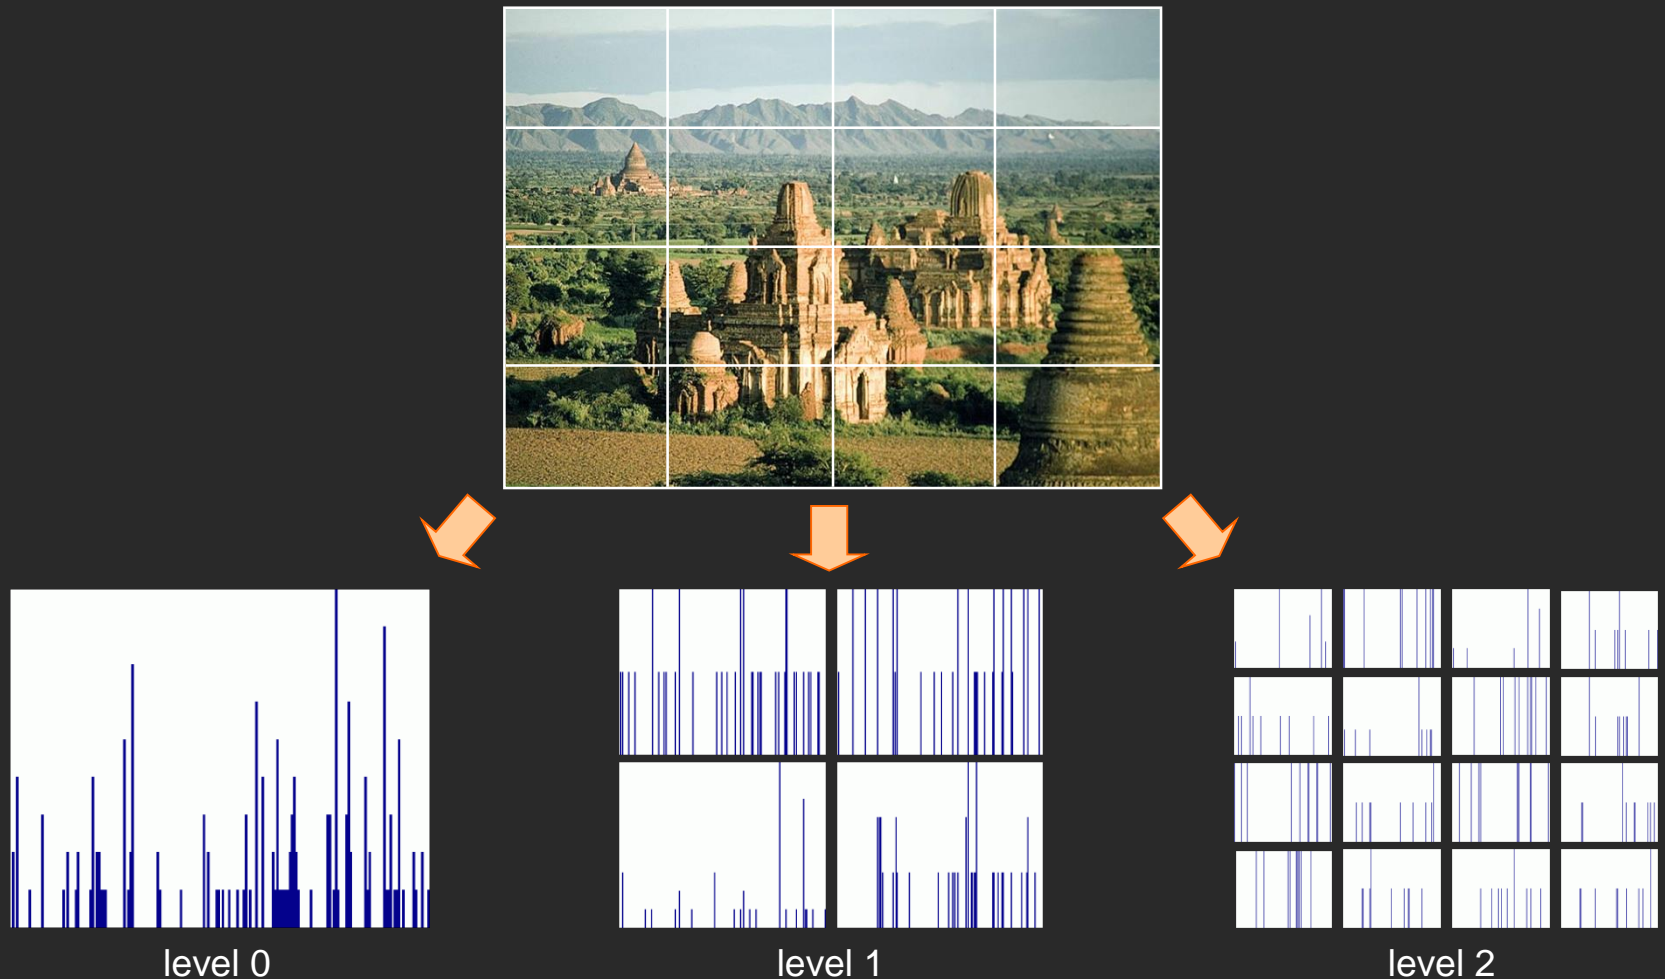
# SIFT vector formation

- 4x4 array of gradient orientation histogram weighted by magnitude
- 8 orientations x 4x4 array = 128 dimensions
- Motivation: some sensitivity to spatial layout, but not too much.



# Spatial pyramid representation

- Extension of a bag of features
- Locally orderless representation at several levels of resolution



# Recap – Image Classification with Bags of Local Features

- Food for thought, doesn't the spatial pyramid seem kind of recursive / hierarchical? Like a SIFT feature on top of SIFT features?
- Seems like there is a tendency for features to involve convolution, spatial pooling, and non-linearities.

# Object Detection

- Overview
- Viola-Jones
- Dalal-Triggs
- Later classes:
  - Deformable models
  - Deep learning

# Person detection with HoG's & linear SVM's



- Histograms of Oriented Gradients for Human Detection, [Navneet Dalal](#), [Bill Triggs](#), International Conference on Computer Vision & Pattern Recognition - June 2005
- <http://lear.inrialpes.fr/pubs/2005/DT05/>


# Object detection vs Scene Recognition

- What's the difference?
- Objects (even if deformable and articulated) probably have more consistent shapes than scenes.
- Scenes can be defined by distribution of “stuff” – materials and surfaces with arbitrary shape.
- Objects are “things” that own their boundaries
- Bag of words models are less popular for object detection because they throw away shape info.



# Object Category Detection

- Focus on object search: “Where is it?”
- Build templates that quickly differentiate object patch from background patch



# Challenges in modeling the object class



Illumination



Object pose



Clutter



Occlusions



Intra-class  
appearance




Viewpoint

# Challenges in modeling the non-object class


True  
Detections




Bad  
Localization



Confused with  
Similar Object



Misc. Background

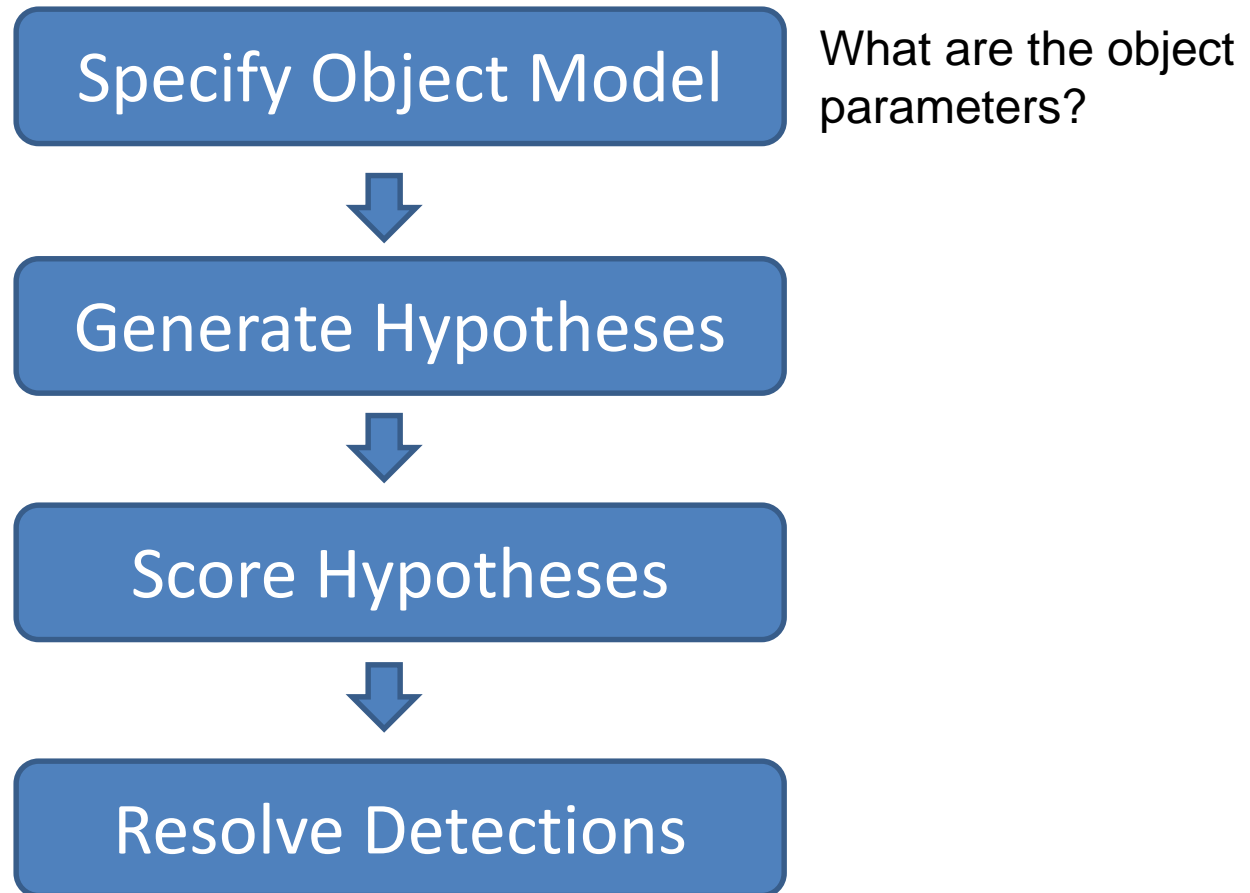


Confused with  
Dissimilar Objects





# General Process of Object Recognition




# Specifying an object model

## 1. Statistical Template in Bounding Box

- Object is some  $(x,y,w,h)$  in image
- Features defined wrt bounding box coordinates



Image





Template Visualization

# Specifying an object model

## 2. Articulated parts model








- Object is configuration of parts
- Each part is detectable



# Specifying an object model

## 3. Hybrid template/parts model

Detections



Template Visualization

root filters  
coarse resolution


part filters  
finer resolution

deformation  
models

# Specifying an object model

## 4. 3D-ish model

- Object is collection of 3D planar patches under affine transformation





# General Process of Object Recognition

Specify Object Model



Generate Hypotheses

Propose an alignment of the model to the image



Score Hypotheses



Resolve Detections

# Generating hypotheses

## 1. Sliding window

- Test patch at each location and scale



# Generating hypotheses

## 1. Sliding window

- Test patch at each location and scale



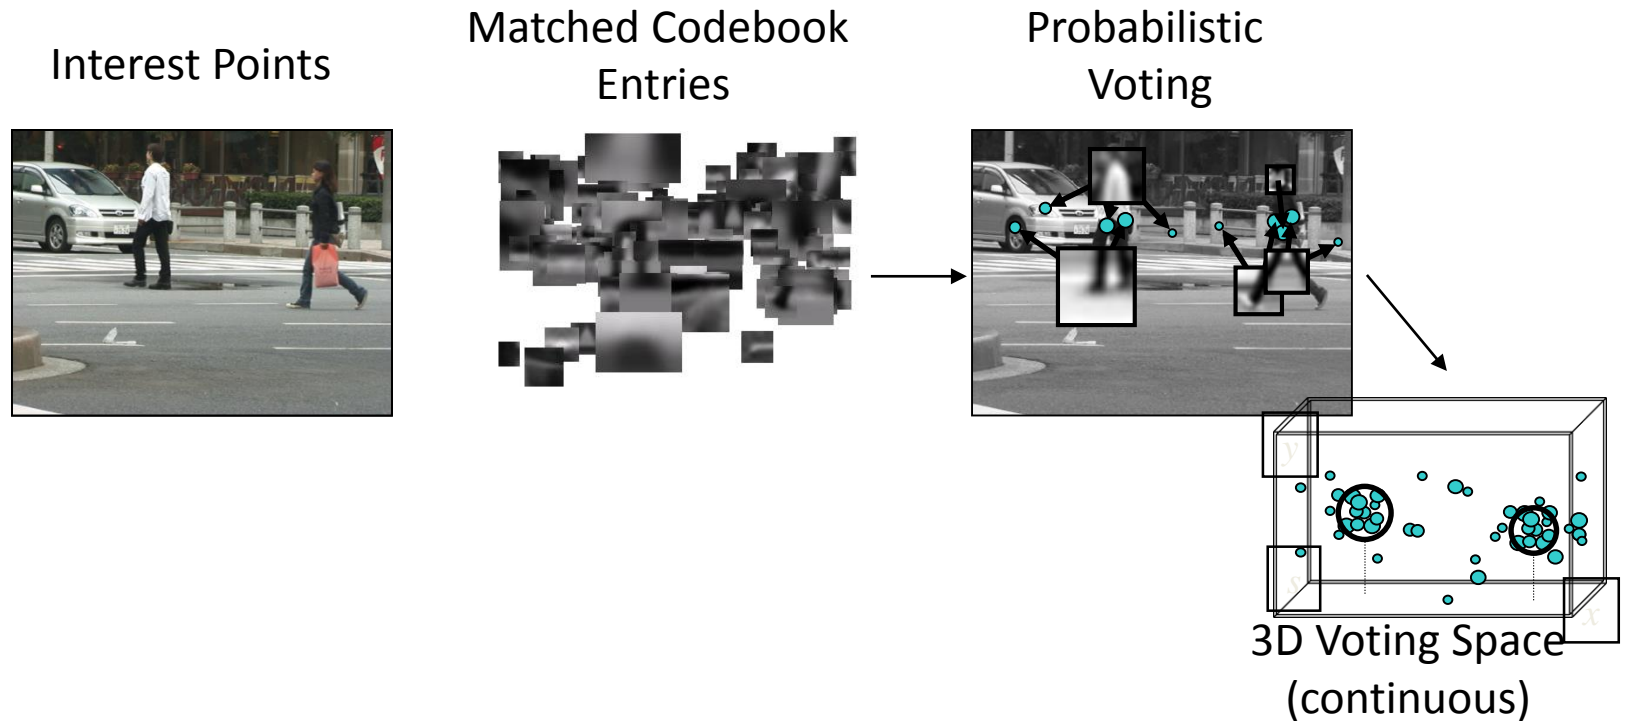
Note – Template did not change size

# Each window is separately classified



# Generating hypotheses






## 2. Voting from patches/keypoints





# Generating hypotheses

## 3. Region-based proposal



# General Process of Object Recognition

Specify Object Model



Generate Hypotheses



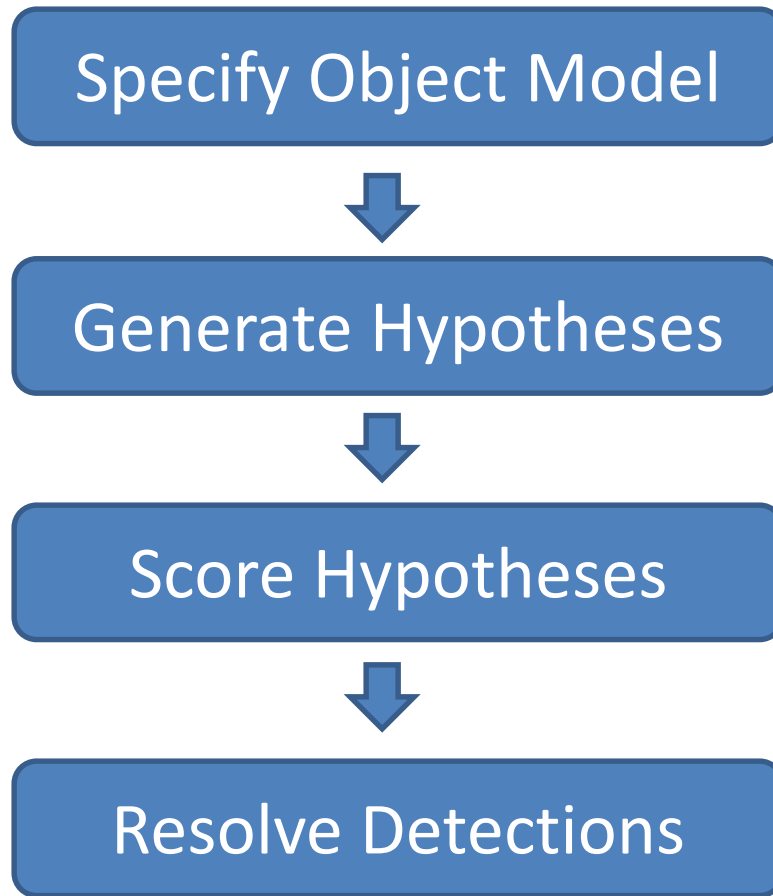
Score Hypotheses

Mainly-gradient based features, usually based on summary representation, many classifiers



Resolve Detections

# General Process of Object Recognition

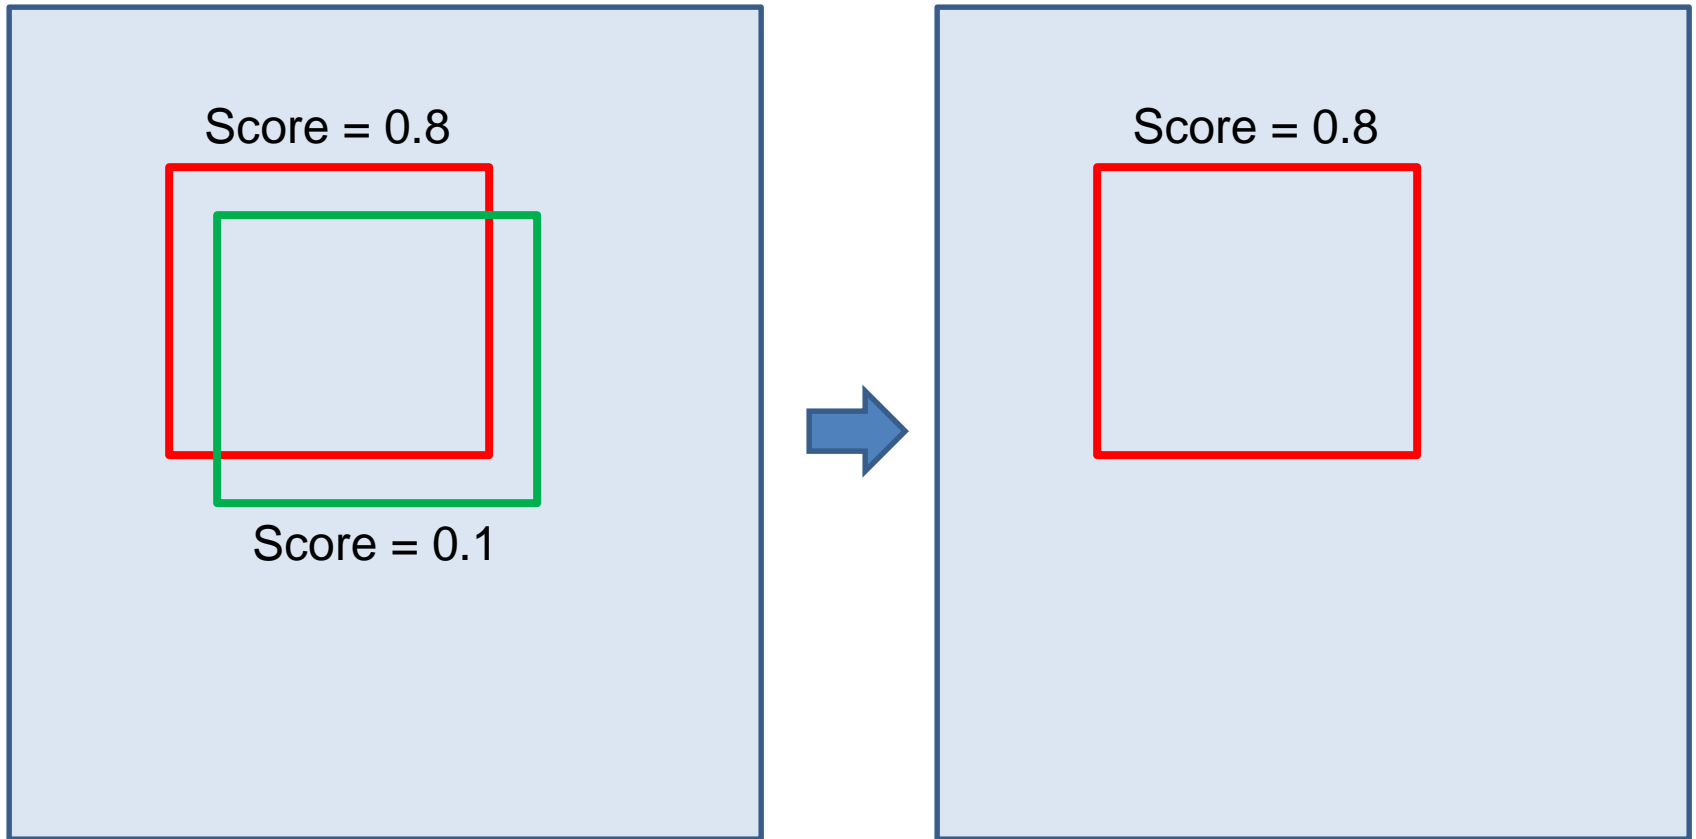


Rescore each proposed object based on whole set



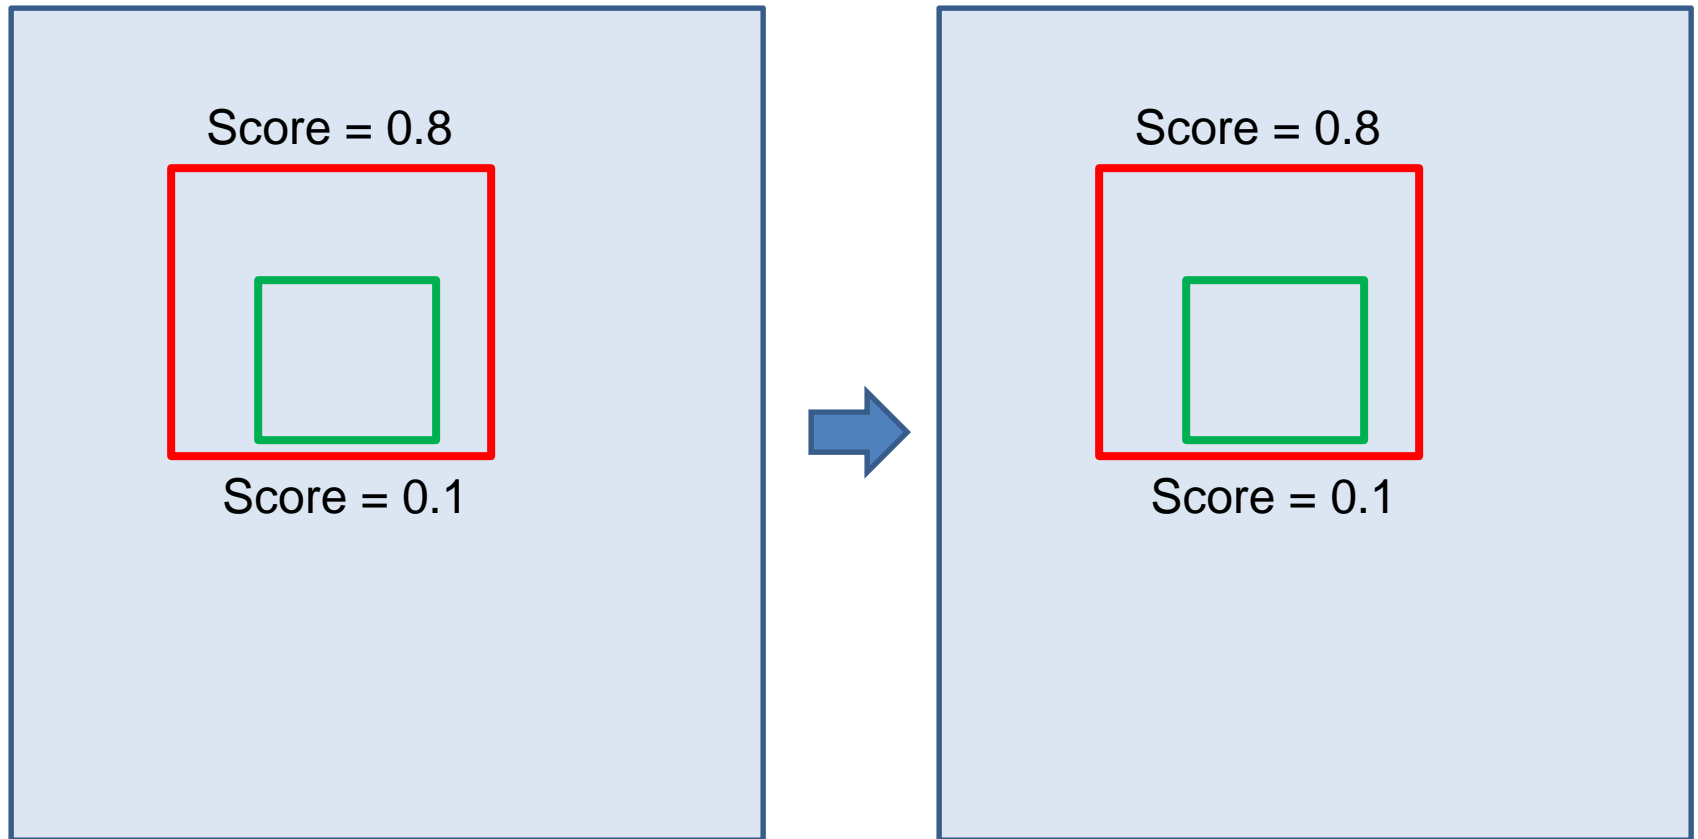
# Resolving detection scores

## 1. Non-max suppression



# Resolving detection scores


## 1. Non-max suppression



“Overlap” score is below some threshold

# Resolving detection scores


## 2. Context/reasoning



(g) Car Detections: Local



(h) Ped Detections: Local



# Influential Works in Detection

- Sung-Poggio (1994, 1998) : ~2000 citations
  - Basic idea of statistical template detection, bootstrapping to get “face-like” negative examples, multiple whole-face prototypes (in 1994)
- Rowley-Baluja-Kanade (1996-1998) : ~3600 citations
  - “Parts” at fixed position, non-maxima suppression, simple cascade, rotation, pretty good accuracy, fast
- Schneiderman-Kanade (1998-2000,2004) : ~1700 citations
  - Careful feature engineering, excellent results, cascade
- Viola-Jones (2001, 2004) : ~16,000 citations
  - Haar-like features, Adaboost as feature selection, hyper-cascade, very fast
- Dalal-Triggs (2005) : ~20,000 citations
  - Careful feature engineering, excellent results, HOG feature, easy to implement
- Felzenszwalb-McAllester-Ramanan (2008): ~4,600 citations
  - Template/parts-based blend
- Girshick et al. (2013): ~2600 citations
  - R-CNN / Fast R-CNN / Faster R-CNN. Deep learned models on object proposals.



# Sliding Window Face Detection with Viola-Jones

# Face detection and recognition

---



Detection



Recognition

"Sally"



# Consumer application: Apple iPhoto

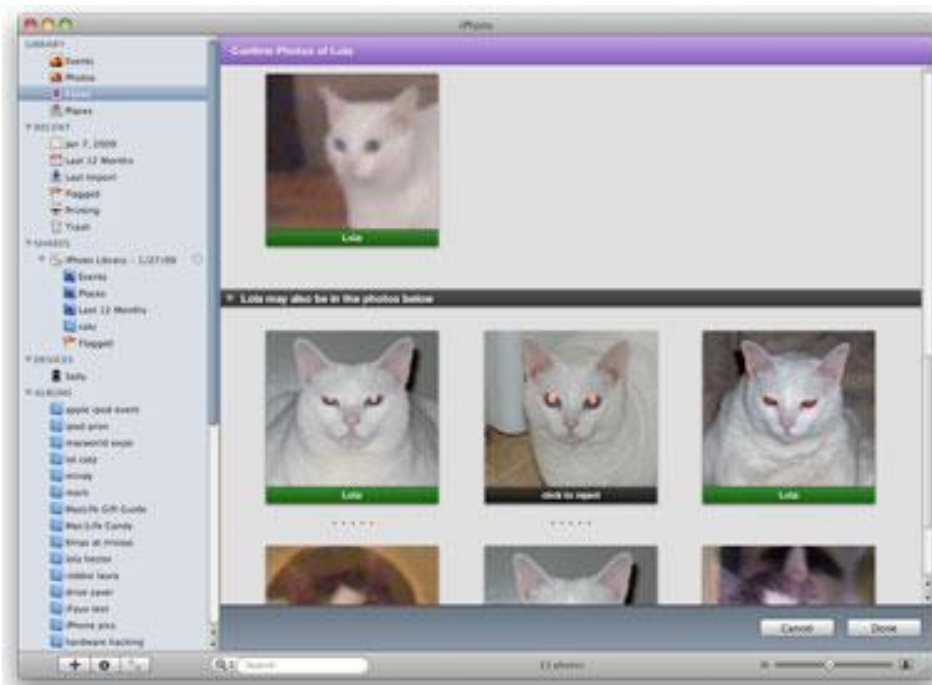
---



<http://www.apple.com/ilife/iphoto/>

# Consumer application: Apple iPhoto

Can be trained to recognize pets!



[http://www.maclife.com/article/news/iphotos\\_faces\\_recognizes\\_cats](http://www.maclife.com/article/news/iphotos_faces_recognizes_cats)



# Consumer application: Apple iPhoto

---

## Things iPhoto thinks are faces



# Funny Nikon ads

---

**"The Nikon S60 detects up to 12 faces."**






# Funny Nikon ads

---

**"The Nikon S60 detects up to 12 faces."**



The Nikon S60. Detects up to 12 faces.

# Challenges of face detection

---

- Sliding window detector must evaluate tens of thousands of location/scale combinations
- Faces are rare: 0–10 per image
  - For computational efficiency, we should try to spend as little time as possible on the non-face windows
  - A megapixel image has  $\sim 10^6$  pixels and a comparable number of candidate face locations
  - To avoid having a false positive in every image image, our false positive rate has to be less than  $10^{-6}$

# The Viola/Jones Face Detector

---

- A seminal approach to real-time object detection
- Training is slow, but detection is very fast
- Key ideas
  - *Integral images* for fast feature evaluation
  - *Boosting* for feature selection
  - *Attentional cascade* for fast rejection of non-face windows



P. Viola and M. Jones. [Rapid object detection using a boosted cascade of simple features.](#) CVPR 2001.

P. Viola and M. Jones. [Robust real-time face detection.](#) IJCV 57(2), 2004.

# Image Features

---

“Rectangle filters”




*Value =*

$$\sum (\text{pixels in white area}) - \sum (\text{pixels in black area})$$

# Example


---



Source




Result



# Fast computation with integral images

---


- The *integral image* computes a value at each pixel  $(x,y)$  that is the sum of the pixel values above and to the left of  $(x,y)$ , inclusive
- This can quickly be computed in one pass through the image






# Computing the integral image

---



# Computing the integral image

---



Cumulative row sum:  $s(x, y) = s(x-1, y) + i(x, y)$

Integral image:  $ii(x, y) = ii(x, y-1) + s(x, y)$

MATLAB: `ii = cumsum(cumsum(double(i)), 2);`


# Computing sum within a rectangle

---

- Let A,B,C,D be the values of the integral image at the corners of a rectangle
- Then the sum of original image values within the rectangle can be computed as:


$$\text{sum} = A - B - C + D$$

- Only 3 additions are required for any size of rectangle!



# Computing a rectangle feature


---



# Feature selection

---

- For a 24x24 detection region, the number of possible rectangle features is ~160,000!



# Feature selection

---

- For a 24x24 detection region, the number of possible rectangle features is ~160,000!
- At test time, it is impractical to evaluate the entire feature set
- Can we create a good classifier using just a small subset of all possible features?
- How to select such a subset?



# Boosting

---

- *Boosting* is a learning scheme that combines *weak learners* into a more accurate *ensemble classifier*
- Weak learners based on rectangle filters:

$$h_t(x) = \begin{cases} 1 & \text{if } p_t f_t(x) > p_t \theta_t \\ 0 & \text{otherwise} \end{cases}$$

Diagram annotations for the weak classifier equation:

- $h_t(x)$ : window
- $f_t(x)$ : value of rectangle feature
- $p_t$ : parity
- $\theta_t$ : threshold

- Ensemble classification function:

$$C(x) = \begin{cases} 1 & \text{if } \sum_{t=1}^T \alpha_t h_t(x) > \frac{1}{2} \sum_{t=1}^T \alpha_t \\ 0 & \text{otherwise} \end{cases}$$

Diagram annotations for the ensemble classification function:

- $\alpha_t$ : learned weights

# Training procedure

---


- Initially, weight each training example equally
- In each boosting round:
  - Find the weak learner that achieves the lowest *weighted* training error
  - Raise the weights of training examples misclassified by current weak learner
- Compute final classifier as linear combination of all weak learners (weight of each learner is directly proportional to its accuracy)
  - Exact formulas for re-weighting and combining weak learners depend on the particular boosting scheme (e.g., AdaBoost)

Y. Freund and R. Schapire, [A short introduction to boosting](#), *Journal of Japanese Society for Artificial Intelligence*, 14(5):771-780, September, 1999.

# Boosting intuition


---

**Weak  
Classifier 1**




# Boosting illustration

---



# Boosting illustration


---



# Boosting illustration

---


**Weights  
Increased**





# Boosting illustration


---



# Boosting illustration

---


**Final classifier is  
a combination of weak  
classifiers**



# Boosting for face detection

---

- First two features selected by boosting:




This feature combination can yield 100% recall and 50% false positive rate

# Boosting for face detection

---

- A 200-feature classifier can yield 95% detection rate and a false positive rate of 1 in 14084




Receiver operating characteristic (ROC) curve

# Attentional cascade


---

- We start with simple classifiers which reject many of the negative sub-windows while detecting almost all positive sub-windows
- Positive response from the first classifier triggers the evaluation of a second (more complex) classifier, and so on
- A negative outcome at any point leads to the immediate rejection of the sub-window





# Attentional cascade

- Chain classifiers that are progressively more complex and have lower false positive rates:



Receiver operating characteristic






# Attentional cascade

---

- The detection rate and the false positive rate of the cascade are found by multiplying the respective rates of the individual stages
- A detection rate of 0.9 and a false positive rate on the order of  $10^{-6}$  can be achieved by a 10-stage cascade if each stage has a detection rate of 0.99 ( $0.99^{10} \approx 0.9$ ) and a false positive rate of about 0.30 ( $0.3^{10} \approx 6 \times 10^{-6}$ )



# Training the cascade

---

- Set target detection and false positive rates for each stage
- Keep adding features to the current stage until its target rates have been met
  - Need to lower AdaBoost threshold to maximize detection (as opposed to minimizing total classification error)
  - Test on a *validation set*
- If the overall false positive rate is not low enough, then add another stage
- Use false positives from current stage as the negative training examples for the next stage

# The implemented system

---

- Training Data
  - 5000 faces
    - All frontal, rescaled to 24x24 pixels
  - 300 million non-faces
    - 9500 non-face images
  - Faces are normalized
    - Scale, translation
- Many variations
  - Across individuals
  - Illumination
  - Pose







# System performance

---

- Training time: “weeks” on 466 MHz Sun workstation
- 38 layers, total of 6061 features
- Average of 10 features evaluated per window on test set
- “On a 700 Mhz Pentium III processor, the face detector can process a 384 by 288 pixel image in about .067 seconds”
  - 15 Hz
  - 15 times faster than previous detector of comparable accuracy (Rowley et al., 1998)

# Output of Face Detector on Test Images

---




# Other detection tasks

---




Facial Feature Localization



Profile Detection


Male vs.  
female






# Profile Detection

---



# Profile Features

---



# Summary: Viola/Jones detector

---

- Rectangle features
- Integral images for fast computation
- Boosting for feature selection
- Attentional cascade for fast rejection of negative windows